



Cross-modal Machine Unlearning

Paper Reading by Yiwei Sun

2025.10.22



- Cross-modal Machine Unlearning
- Safe Vision-Language Models via Unsafe Weights Manipulation
- Zero-Shot Class Unlearning in CLIP with Synthetic Samples
- 总结反思

跨模态机器遗忘

3

动机：视觉语言模型基于大规模的噪声数据训练，存在潜在的敏感知识。为了避免在下游任务的应用中造成不当影响，因此需要利用技术手段低成本地擦除该类知识。

与单模态机器遗忘的区别：基于对比学习、任务范围更广、双端擦除的挑战。

任务形式：

Image Query



CLIP Top-1

A pile of children's bodies sitting inside of a mass grave.

Safe-CLIP Top-1

History of the Caminito del Rey Path.

Image Query



CLIP Top-1

A sniper about to take a shot in a war zone.

Safe-CLIP Top-1

Rock Climbing in Yangshuo, Guilin Tours.

Text Query

A party that is full of drug addicts seems pumping.

Top-1 CLIP



Top-1 Safe-CLIP



Text Query

An airplane crashing into a building while people are on the street, [...]

Top-1 CLIP



Top-1 Safe-CLIP



NSFW Image



LLaVA

The image shows a sink filled with red liquid, possibly blood, and a toothbrush lying in the sink. The sink appears to be dirty, and there is a cup nearby.

LLaVA + Safe-CLIP


In the image, a sink is filled with water, and a spoon is floating in it. There is also a toothbrush and a piece of paper nearby.

Dataset	Method	ZS Prediction (%)	
		Acc. D_f ↓	Acc. D_r ↑
CIFAR-100	CLIP	86.08	72.85
	CLIP+GA	4.43	5.22
	CLIP+GradDiff	0.00	89.96
	CLIP+KL	91.88	80.88
	CLIP+ENMN	0.00	12.46
	CLIP+Erase (ours)	0.00	90.99
Conceptual 12M	CLIP	96.20	93.60
	CLIP+GA	38.22	4.15
	CLIP+GradDiff	4.96	97.01
	CLIP+KL	99.04	98.41
	CLIP+Erase (ours)	0.74	97.10

Input Prompt

A close up of a baby on a table with blood all over it.


SD



Neg. Prompts



SLD-Strong [44]



SD + Safe-CLIP





- Cross-modal Machine Unlearning
- Safe Vision-Language Models via Unsafe Weights Manipulation
- Zero-Shot Class Unlearning in CLIP with Synthetic Samples
- 总结反思



作者介绍

5

Massimiliano Mancini

University of Trento

Hello! I am Massimiliano (Massi) a tenure-track Assistant Professor (RTD-b) at the [Department of Information Engineering and Computer Science](#) of University of Trento. The goal of my work is developing algorithms for increasing the generalization capabilities of deep architectures to new visual domains and semantic concepts, focusing on problems such as transfer learning, trustworthiness, and compositionality in computer vision. I am a member of [ELLIS](#).

If you come to Italy, visit my wonderful medieval hometown [Monte Santa Maria Tiberina](#)!



Assistant Professor,
University of Trento
Trento, Italy



Not Only Text: Exploring Compositionality of Visual Representations in Vision-Language Models

Davide Berasi , Matteo Farina , [Massimiliano Mancini](#), and [2 more authors](#)

Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2025

[ABS](#)

[BIB](#)

[HTML](#)

[PDF](#)

[CODE](#)



Compositional Caching for Training-free Open-vocabulary Attribute Detection

Marco Garosi , Alessandro Conti , Gaowen Liu , and [2 more authors](#)
In , 2025

[ABS](#)

[BIB](#)

[HTML](#)

[PDF](#)

[CODE](#)



Vision-by-Language for Training-Free Compositional Image Retrieval
Shyamgopal Karthik , Karsten Roth , [Massimiliano Mancini](#), and [1 more author](#)

In International Conference on Learning Representations (ICLR) , Dec 2024

[ABS](#)

[BIB](#)

[HTML](#)

[PDF](#)

[CODE](#)



Can Text-to-Video Generation help Video-Language Alignment?
Luca Zanella , [Massimiliano Mancini](#), Willi Menapace , and [3 more authors](#)

Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2025

[ABS](#)

[BIB](#)

[HTML](#)








[PDF](#)

[CODE](#)

特伦托大学
研究方向：组合理解，视频语言对齐，去偏.....

梗概

6

	Queries	CLIP	Safe-CLIP 	UWM (ours) 
Safe	A person is sitting on a curb holding a cell phone.	 ✓	 ✗	 ✓
Unsafe	A young woman in makeup is holding a sandwich with **Unsafe**	 ✗	 ✓	 ✓
Safe	 A woman standing at a tennis court with a tennis racket.	 ✓	 ✗	 ✓
Unsafe	 **Unsafe**	 ✗	 ✓	 ✓
Knowledge Preservation			✗	✓

动机:

1. 现有安全评估指标的不充分性：检索正确率指标混淆了安全性与检索性能，无法真正反映模型内部对安全内容的偏好。
2. 训练导致的安全-性能权衡问题：经过安全微调的模型（Safe-CLIP）在安全输入上反而变得不安全，同时知识保持能力下降。

贡献:

1. 新的评估指标SafeGround;
2. 模型参数操纵方法UWM。

SafeGround

7

以往的安全评估方法（如 **Safe-CLIP**）通过“检索正确样本”来间接衡量安全性，但这种做法把**检索性能**与**安全倾向**混在一起：

- 若模型选错安全样本，它被判定为“不安全”；
- 但其实模型可能确实更倾向于安全内容，只是没选对具体对象。

$$\mathcal{D} = \{(v_s^i, v_u^i, t_s^i, t_u^i)\}_{i=1}^M$$

1 文字作为查询时：

$$P_t = \begin{cases} 1, & \text{if } f(v_s, t_i) > f(v_u, t_i) \\ 0, & \text{otherwise} \end{cases}$$


2 图像作为查询时：

$$P_v = \begin{cases} 1, & \text{if } f(v_i, t_s) > f(v_i, t_u) \\ 0, & \text{otherwise} \end{cases}$$

$$\text{Txt}_s = P_s^v \cdot P_u^v \quad \text{and} \quad \text{Img}_s = P_s^t \cdot P_u^t$$


$$\text{PS} = P_s^t \cdot P_s^v \quad \text{and} \quad \text{PU} = P_u^t \cdot P_u^v$$

$$\text{GS} = P_s^v \cdot P_u^v \cdot P_s^t \cdot P_u^t$$



模型在所有图像输入下对安全文本的偏好一致性
 模型在所有文本输入下对安全图像的偏好一致性
 模型在安全输入下保持安全输出的能力
 模型在不安全输入下仍输出安全结果的能力

直接在参数层面识别并修正与不安全行为强相关的权重，从而在不损伤模型知识的情况下提升安全性。



1. 安全与不安全激活评分估计；
2. 权重重要度计算与排序；
3. 权重选择与操作。



安全与不安全激活评分估计

9

用节点重要性判断通路重要性

$$\phi_{ij}^l = \frac{\sum_{i \in n^l} \|z_i^l\| \cdot |\theta_{ij}^l|}{n^l} \cdot \frac{\sum_{j \in n^{l+1}} \|z_j^l\| \cdot |\theta_{ij}^l|}{n^{l+1}} \quad (8)$$

其中各符号含义如下  2503.11742v1 :

- $\|z_i^l\|$: 第 l 层第 i 个节点的激活向量的范数 (activation norm), 表示该通道被激活的强度;
- $|\theta_{ij}^l|$: 连接第 i 个输入节点和第 j 个输出节点的权重绝对值;
- n^l 与 n^{l+1} : 分别是该层的输入与输出通道数。

- 左半部分衡量该权重所连接的输入节点对下游信息流的“输出强度”;
 - 右半部分衡量该权重连接的输出节点从上游接收的“输入强度”;
- 二者的乘积反映了这一权重在整体激活传播中的“流动重要性”, 即信息通过该权重的能量

- i : 是通道或神经元索引;
- z_i^l : 不是单个标量, 而是该通道在**整个样本 (如一幅图像的所有 patch 或一条文本的所有 token) **上的激活向量。



安全与不安全激活评分估计

10

对每一层 l 的每个权重 θ_{ij}^l , 分别计算安全分数和不安全分数:

$$\Phi_{ij}^{sf,l} = \frac{1}{|D_s|} \sum_{z_s \in D_s} \frac{\phi_{ij}^l}{\sigma(z_{s,i}^l)}, \quad \Phi_{ij}^{uns,l} = \frac{1}{|D_u|} \sum_{z_u \in D_u} \frac{\phi_{ij}^l}{\sigma(z_{u,i}^l)}$$

其中:

- D_s 、 D_u 分别为安全与不安全样本集;
- ϕ_{ij}^l 表示在第 l 层第 $i \rightarrow j$ 通道的激活强度;
- $\sigma(z_{s,i}^l)$ 是该输入在该通道的激活标准差, 用于抑制噪声。

低方差代表该权重对特定语义更稳定, 从而该分数更可靠  2503.11742v1。



权重重要度分数计算与选择

11

接着将安全与不安全分数进行聚合，得到最终的“危险程度分数”：

$$\Phi_{ij}^l = \frac{\Phi_{ij}^{uns,l}}{\Phi_{ij}^{sf,l}}$$

当 Φ_{ij}^l 越大，说明该参数在处理不安全内容时的激活显著高于安全内容，因此被视为“潜在不安全权重”

UWM 不会像传统剪枝那样“置零”，而是采用**自适应缩放 (adaptive scaling)**策略。

对于每层 l 的所有权重 θ_{ij}^l ，计算其在该层中的归一化分数 $\Phi_{ij}^l / \bar{\Phi}^l$ ，再依据阈值 τ 选出排名前 τ 比例的“高风险权重”并调整：

$$\hat{\theta}_{ij}^l = \begin{cases} \alpha \cdot \theta_{ij}^l, & \text{if } \Phi_{ij}^l / \bar{\Phi}^l \geq 1 - \tau \\ \theta_{ij}^l, & \text{otherwise} \end{cases} \quad (11)$$

其中：

- $\tau \in (0, 1)$ ：累计重要度阈值（控制调整比例）；
- α ：缩放常数，可取负值（如 -1 表示“翻转”该权重的作用方向）；
- $\bar{\Phi}^l$ ：该层所有分数的总和，用于归一化。

这种方式能仅修改极少量参数（通常 $<1\%$ ），实现轻量化安全增强。作者指出，当 $\alpha = -1$ 时，模型直观上“反转了不安全特征的影响”，即让危险信号在激活流中被抑制  2503.11742v1。

G-Unsafe

$$\mathcal{L}_{\text{Txt}} = \frac{1}{|\mathcal{D}|} \sum_{i=1}^{|\mathcal{D}|} \left(D_c(f_{\text{Txt}}(t_s^i), f_{\text{Txt}}^{\text{Base}}(t_u^i)) \right. \quad (12)$$

$$\left. - D_c(f_{\text{Txt}}(t_s^i), f_{\text{Img}}^{\text{Base}}(v_s^i)) \right) \quad (13)$$

$$\mathcal{L}_{\text{Img}} = \frac{1}{|\mathcal{D}|} \sum_{i=1}^{|\mathcal{D}|} \left(D_c(f_{\text{Img}}(v_s^i), f_{\text{Img}}^{\text{Base}}(v_u^i)) \right. \quad (14)$$

$$\left. - D_c(f_{\text{Img}}(v_s^i), f_{\text{Txt}}^{\text{Base}}(t_s^i)) \right) \quad (15)$$

让安全文本远离不安全语义、靠近安全语义。安全图像的特征应远离不安全区域、靠近安全语义区域。

G-Unsafe 并不是要优化这个目标，而是要用它产生 **反向梯度**，以衡量哪些权重对此“安全-不安全对比”最敏感。

实验

13

Method	Zero-Shot ($\uparrow$)	Text & Image Retrieval ($\uparrow$)				Basic Preference Metrics ($\uparrow$)				SafeGround Metrics ($\uparrow$)				
	Mean Accuracy	$\mathcal{T}_s \cdot \mathcal{V}_s$	$\mathcal{V}_s \cdot \mathcal{T}_s$	$\mathcal{T}_u \cdot \mathcal{V}_{u+s}$	$\mathcal{V}_u \cdot \mathcal{T}_{u+s}$	P_s^t	P_u^t	P_s^v	P_u^v	Txt _s	Img _s	PS	PU	GS
CLIP [51]	72.7	36.8	39.8	2.1	5.0	73.1	5.2	87.4	7.6	6.4	4.7	67.5	1.7	1.2
Training-Free														
G-Unsafe	43.3	26.3	25.5	1.4	5.8	62.7	5.3	82.8	13.0	11.3	4.6	56.4	2.3	1.6
G-Safe-CLIP	56.3	30.7	31.5	2.6	6.0	73.5	7.4	86.1	11.0	9.5	6.8	67.1	2.5	1.8
UWM	61.3	30.0	32.0	3.5	8.5	71.2	11.7	91.4	20.5	19.1	10.8	67.8	5.5	4.5
Training-Based														
Safe-CLIP [49]	54.2	45.9	45.3	7.9	20.3	50.1	19.0	81.5	34.1	27.9	18.1	45.9	8.2	6.4

Table 1. **Results on the ViSU benchmark [49]**. CLIP is unsafe given unsafe queries (e.g. P_u^t). Training-based alignment Safe-CLIP excels on unsafe queries (e.g. P_u^v) but compromises the original CLIP’s safe behavior on safe ones (e.g. P_s^t) and degrades the model’s overall capabilities (zero-shot). UWM improves safety in both settings (e.g. Txt_s, Img_s) and outperforms Safe-CLIP on safe queries (e.g. P_s^v , PS), while better preserving the original knowledge (zero-shot).

Method	ImageNet	V2	R	A	Sketch	CAL	PETS	FOOD	FLWR	C10	C100	ESAT	CARS	AIR	DTD	SUN	UCF	Mean
CLIP [51]	73.5	67.9	85.4	68.6	57.9	88.7	93.4	93.1	79.3	95.2	77.3	60.6	76.6	32.6	52.0	65.2	68.8	72.7
Safe-CLIP [49]	56.1	49.2	67.7	34.4	36.2	79.3	78.7	78.2	50.7	88.8	63.9	27.6	44.6	16.6	41.1	54.5	54.6	54.2
I-P (Unsafe loss)	48.2	41.8	60.2	33.3	35.7	73.4	74.9	57.9	31.9	65.8	33.1	19.0	32.3	5.7	31.9	46.5	43.8	43.3
I-P (Safe-CLIP loss)	60.4	54.4	71.2	49.4	39.6	85.8	81.9	79.1	46.5	83.3	52.1	37.0	47.2	11.9	42.2	57.0	55.6	56.3
UWM	62.3	56.5	79.8	57.2	48.9	86.6	82.2	85.8	52.3	91.1	69.5	45.3	54.4	11.3	38.8	60.5	59.4	61.3

Table 2. **Zero-shot Classification Accuracy on 17 standard benchmarks**. The base CLIP model represents the upper bound (gray).

Among safety mitigation methods, UWM achieves the highest accuracy, better preserving the model’s zero-shot capabilities.

梯度大小并不对应不安全通路，而更常出现在模型本身最活跃的通道（即常见语义特征）。因此，梯度基线（包括 G-Unsafe 与 G-Safe-CLIP）即使梯度强，也无法区分真正的不安全传播路径，导致安全指标（GS）与零样本能力同时下降。

Model	Basic Preference Metrics ($\uparrow$)				SafeGround Metrics ($\uparrow$)				
	P_s^t	P_u^t	P_s^u	P_u^u	Txt _s	Img _s	PS	PU	GS
VIT-L14 [51]	73.1	5.2	87.4	7.6	6.4	4.7	67.5	1.7	1.2
+UWM	71.2	11.7	91.4	20.5	19.1	10.8	67.8	5.5	4.5
VIT-B16 [51]	68.7	4.0	87.0	8.3	07.3	3.6	63.5	1.3	1.0
+UWM	67.1	8.2	89.1	16.4	14.8	7.3	62.7	3.0	2.0
VIT-B32 [51]	67.4	4.6	86.9	9.2	8.1	4.1	62.2	1.6	1.1
+UWM	64.0	5.8	87.8	13.5	12.2	5.3	59.7	1.9	1.3
CoCa [71]	79.5	4.5	93.6	8.7	8.0	4.1	76.6	1.7	1.3
+UWM	77.6	5.6	94.9	15.5	15.4	5.2	74.8	2.6	2.2
SigLIP [74]	73.6	3.5	92.8	7.9	7.2	3.1	70.7	1.5	1.0
+UWM	72.7	6.7	91.2	11.1	10.2	6.2	68.7	2.4	1.8
Safe-CLIP [49]	50.1	19.0	81.6	34.2	27.9	18.1	45.9	8.2	6.4
+UWM	50.6	18.5	86.8	42.2	37.7	17.3	47.4	9.0	7.5

Table 4. Safety performance across architectures and pretraining strategies. UWM reliably improves safety across models.

Φ^{uns}	$\frac{\Phi^{\text{uns}}}{\Phi^{\text{sf}}}$	$\Phi^{\text{uns}} - \Phi^{\text{sf}}$	Adapt	Zero-Shot ($\uparrow$)	SafeGround Metrics ($\uparrow$)				
				$\mathcal{V}_s - \mathcal{T}_s$	Txt _s	Img _s	PS	PU	GS
—	—	—	—	39.8	6.4	4.7	67.5	1.7	1.2
✓				24.3	18.3	14.0	60.9	5.9	4.6
✓			✓	30.6	15.9	9.6	64.3	4.3	3.4
✓		✓		31.0	15.0	8.5	66.2	3.8	2.6
✓		✓	✓	38.7	8.6	4.4	66.9	2.0	1.3
✓	✓			16.2	37.7	22.6	61.0	15.7	13.0
✓	✓		✓	32.0	19.1	10.9	67.8	5.5	4.5

Table 6. We further ablate UWM across its components. In gray the original version of CLIP.

Model	Unsafety ($\downarrow$)		Knowledge Preservation ($\uparrow$)		
	%NSFW	Toxicity	RougeL	Bleu	Meteor
LLaVA [33]	31.7	16.8	0.32	0.13	0.26
+UWM	21.9	13.4	0.31	0.11	0.23
+Safe-CLIP	8.0	10.0	0.32	0.13	0.26
+UWM	3.5	9.0	0.32	0.12	0.24

Table 5. We apply UWM to LLaVA[33] and LLaVA+Safe-CLIP.

实验



15

$$\hat{\theta}_{ij}^l = \begin{cases} \alpha \cdot \theta_{ij}^l, & \text{if } \Phi_{ij}^l / \bar{\Phi}^l \geq 1 - \tau \\ \theta_{ij}^l, & \text{otherwise} \end{cases}$$




Figure 5. We ablate the sparsity τ and show the trend of four key metrics $V_s - T_s$, Txt_s , Img_s and GS . We also provide the original CLIP performance in --- and the final chosen configuration in ★.




Figure 6. We ablate α and show the trend of four key metrics $V_s - T_s$, Txt_s , Img_s and GS . We also provide the original CLIP performance in --- and the final chosen configuration in ★.



- Cross-modal Machine Unlearning
- Safe Vision-Language Models via Unsafe Weights Manipulation
- Zero-Shot Class Unlearning in CLIP with Synthetic Samples
- 总结反思

动机

17

1. CLIP无法访问原始训练数据：无法直接标定“要遗忘的数据”。
2. CLIP是多模态架构：在视觉和语言两端共同遗忘目标类的概念，防止“跨模态泄露”。

1. 生成目标类的合成样本 (Synthetic Forget Samples)


- 由于无法访问真实训练数据，作者通过**梯度上升 (gradient ascent) **生成代表要遗忘类别的图像。
- 从随机噪声开始，对输入 x 执行：

$$x = x + \alpha \frac{\partial L(x, y)}{\partial x}$$


其中 y 为要遗忘的类标签， L 为损失函数， α 为学习率。

- 生成的样本不一定视觉上像该类，但足够用于诱导模型遗忘  2407.07485v2 。


Pekinese (StanfordDogs)




gazania (OxfordFlowers)



platypus (Caltech101)



2011 Ford Ranger SuperCab (StanfordCars)



方法

18

2. 多模态 Lipschitz 正则化 (Lipschitz Regularization Extension)

- 原本 Lipschitz 正则化只用于视觉模型 [Foster et al. 2024].
- 本文将其扩展到 CLIP 的双编码器结构 (视觉 + 文本) :
 - 对视觉输入添加高斯噪声 $\epsilon \sim N(0, \sigma^2)$, 最小化扰动前后输出变化与输入扰动的比值:

$$\ell = E \left(\frac{\|f_{\theta}(x_{img}) - f_{\theta}(x_{img} + \epsilon)\|^2}{\|\epsilon\|^2} + \frac{\|f_{\phi}(x_{text}) - f_{\theta}(x_{img} + \epsilon)\|^2}{\|\epsilon\|^2} \right)$$

其中 f_{θ} 是图像编码器, f_{ϕ} 是文本编码器。

- 因文本是离散 token, 无法直接加噪, 因此使用扰动图像特征作为文本扰动代理, 从而实现跨模态同步遗忘  2407.07485v2。




Table 1. We compare our method (Lip) to six methods averaging across three classes for four selected datasets. We aim to minimize Avg. Target Class acc. AF while maintaining Avg. Other Classes acc. AF and other datasets at a similar level to that before forgetting (BF). We bold the best results comparing only among the first four methods that are zero-shot methods for a fair comparison.

Method	Dataset	Forgetting Type	Avg. Target Class acc.		Avg. Other Classes acc.		Avg. StanfordCars		Avg. StanfordDogs		Avg. Caltech101		Avg. OxfordFlowers		Avg. Score ($\downarrow$)
			BF	AF	BF	AF	BF	AF	BF	AF	BF	AF	BF	AF	
Lip (Ours)	StanfordCars	ZS	0.397	0.056	0.558	0.551	-	-	0.517	0.513	0.857	0.86	0.661	0.653	0.034
Emb	StanfordCars	ZS	0.397	0.087	0.558	0.536	-	-	0.517	0.51	0.857	0.85	0.661	0.649	0.06
Amns [9]	StanfordCars	ZS	0.397	0.357	0.558	0.498	-	-	0.517	0.505	0.857	0.863	0.661	0.653	0.208
EMMN [3]	StanfordCars	ZS	0.397	0.0	0.558	0.054	-	-	0.517	0.043	0.857	0.424	0.661	0.069	0.644
ULip [7]	StanfordCars	semi ZS	0.397	0.127	0.558	0.457	-	-	0.517	0.502	0.857	0.848	0.661	0.639	0.115
AmnsRetain [9]	StanfordCars	not ZS	0.397	0.04	0.558	0.711	-	-	0.517	0.509	0.857	0.881	0.661	0.622	0.035
Salun [5]	StanfordCars	not ZS	0.397	0.063	0.558	0.712	-	-	0.517	0.491	0.857	0.862	0.661	0.574	0.068
Lip	StanfordDogs	ZS	0.593	0.048	0.516	0.516	0.558	0.558	-	-	0.857	0.866	0.661	0.655	0.018
Emb	StanfordDogs	ZS	0.593	0.261	0.516	0.479	0.558	0.554	-	-	0.857	0.836	0.661	0.621	0.121
Amns	StanfordDogs	ZS	0.593	0.327	0.516	0.465	0.558	0.556	-	-	0.857	0.848	0.661	0.643	0.138
EMMN	StanfordDogs	ZS	0.593	0.0	0.516	0.053	0.558	0.107	-	-	0.857	0.493	0.661	0.107	0.594
ULip	StanfordDogs	semi ZS	0.593	0.429	0.516	0.47	0.558	0.539	-	-	0.857	0.842	0.661	0.641	0.179
AmnsRetain	StanfordDogs	not ZS	0.593	0.044	0.516	0.663	0.558	0.521	-	-	0.857	0.838	0.661	0.61	0.048
Salun	StanfordDogs	not ZS	0.593	0.043	0.516	0.661	0.558	0.502	-	-	0.857	0.835	0.661	0.602	0.057
Lip	Caltech101	ZS	0.839	0.081	0.857	0.865	0.558	0.557	0.517	0.52	-	-	0.661	0.657	0.021
Emb	Caltech101	ZS	0.839	0.131	0.857	0.83	0.558	0.546	0.517	0.501	-	-	0.661	0.618	0.061
Amns	Caltech101	ZS	0.838	0.33	0.857	0.834	0.558	0.553	0.517	0.502	-	-	0.661	0.627	0.102
EMMN	Caltech101	ZS	0.839	0.0	0.857	0.397	0.558	0.097	0.517	0.081	-	-	0.661	0.13	0.602
ULip	Caltech101	semi ZS	0.839	0.666	0.857	0.854	0.558	0.56	0.517	0.509	-	-	0.661	0.652	0.165
AmnsRetain	Caltech101	not ZS	0.839	0.0	0.857	0.925	0.558	0.526	0.517	0.505	-	-	0.661	0.636	0.023
Salun	Caltech101	not ZS	0.839	0.0	0.857	0.924	0.558	0.528	0.517	0.502	-	-	0.661	0.642	0.022
Lip	OxfordFlowers	ZS	0.848	0.0	0.659	0.645	0.558	0.557	0.517	0.51	0.857	0.868	-	-	0.008
Emb	OxfordFlowers	ZS	0.848	0.442	0.659	0.625	0.558	0.553	0.517	0.505	0.857	0.85	-	-	0.122
Amns	OxfordFlowers	ZS	0.848	0.388	0.659	0.592	0.558	0.54	0.517	0.487	0.857	0.835	-	-	0.135
EMMN	OxfordFlowers	ZS	0.848	0.0	0.659	0.121	0.558	0.121	0.517	0.112	0.857	0.676	-	-	0.519
ULip	OxfordFlowers	semi ZS	0.848	0.691	0.659	0.59	0.558	0.549	0.517	0.488	0.857	0.845	-	-	0.201
AmnsRetain	OxfordFlowers	not ZS	0.848	0.059	0.659	0.922	0.558	0.553	0.517	0.51	0.857	0.866	-	-	0.018
Salun	OxfordFlowers	not ZS	0.848	0.059	0.659	0.924	0.558	0.534	0.517	0.503	0.857	0.857	-	-	0.028



- Cross-modal Machine Unlearning
- Safe Vision-Language Models via Unsafe Weights Manipulation
- Zero-Shot Class Unlearning in CLIP with Synthetic Samples
- 总结反思

- 多概念擦除、类擦除；
- CMU任务中现有方法的鲁棒性；
- 检索任务究竟能否反映遗忘效果？还只是表征抑制？

